



(a) Cumulative attention ranks windows

re-tiered every 8 decode steps




demoting one window




(b) Cards choose what to read

every decode step, per KV head



(c) One fused pass reads the cache

every decode step, every layer



window mass is added to the cumulative score