


sink
 fp16 window
 local window
 int2 window
 card
 credited from its card
 query / new token
 X dropped


(a) Cumulative attention ranks windows

re-tiered at every window boundary




(b) Cards choose what to read

every decode step, per KV head




cumulative score keeps rolling, read or not



(c) One fused pass per generated token

every decode step, every layer



every window's mass is added to its cumulative score