



■ sink
 ■ fp16 window
 ■ local window
 ■ int2 window
 ■ card
 ■ credited from its card
 ■ query / new token
 ✕ dropped

(a) Cumulative attention ranks windows

re-tiered at every window boundary




demoting one window




(b) Cards choose what to read

every decode step, per KV head




cumulative score keeps rolling, read or not



every window's mass is added to its cumulative score

(c) One fused pass per generated token

every decode step, every layer



credit for unread windows

