



(a) Cumulative attention ranks the windows

Re-tiered at every window boundary




(c) One fused pass per token

Every decode step, every layer



(b) Cards choose what to read

Per KV head, every decode step



Card: read for every window

Opened: its int2 window is read too

Not opened: only its card is read